

Introduction To Programming In R

Last updated November 20, 2013



The Institute
for Quantitative Social Science
at Harvard University

Outline

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Workshop description

- This is an intermediate/advanced R course
 - Appropriate for those with basic knowledge of R
 - Learning objectives:
 - Index data objects by position, name or logical condition
 - Understand looping and branching
 - Write your own simple functions
 - Debug functions
 - Understand and use the S3 object system
-

Running example

Throughout this workshop we will return to a running example that involves calculating descriptive statistics for every column of a `data.frame`. We will often use the built-in *iris* data set. You can load the *iris* data by evaluating `data(iris)` at the R prompt.

Our main example today consists of writing a statistical summary function that calculates the min, mean, median, max, sd, and n for all numeric columns in a `data.frame`, the correlations among these variables, and the counts and proportions for all categorical columns. Typically I will describe a topic and give some generic examples, then ask you to use the technique to start building the summary.

Topic

- 1 Workshop overview and materials
- 2 Data types**
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Vectors and data classes

Values can be combined into vectors using the `c()` function

```
> num.var <- c(1, 2, 3, 4) # numeric vector
> char.var <- c("1", "2", "3", "4") # character vector
> log.var <- c(TRUE, TRUE, FALSE, TRUE) # logical vector
> char.var2 <- c(num.var, char.var) # numbers converted to character
>
```

Vectors have a *class* which determines how functions treat them

```
> class(num.var)
[1] "numeric"
> mean(num.var) # take the mean of a numeric vector
[1] 2.5
> class(char.var)
[1] "character"
> mean(char.var) # cannot average characters
[1] NA
> class(char.var2)
[1] "character"
```

Vector conversion and info

Vectors can be converted from one class to another

```
> class(char.var2)
[1] "character"
> num.var2 <- as.numeric(char.var2) # convert to numeric
> class(num.var2)
[1] "numeric"
> mean(as.numeric(char.var2)) # now we can calculate the mean
[1] 2.5
> as.numeric(c("a", "b", "c")) # cannot convert letters to numeric
[1] NA NA NA
```

In addition to class, you can examine the length() and str() ucture of vectors

```
> ls() # list objects in our workspace
[1] "char.var" "char.var2" "log.var" "num.var" "num.var2"
[6] "tmp"
> length(char.var) # how many elements in char.var?
[1] 4
> str(num.var2) # what is the structure of num.var2?
num [1:8] 1 2 3 4 1 2 3 4
```


Factor vectors

Factors are stored as numbers, but have character labels. Factors are useful for

- Modeling (automatically contrast coded)
- Sorting/presenting values in arbitrary order

Most of the time we can treat factors as though they were character vectors

Lists and data.frames

- A *data.frame* is a list of vectors, each of the same length
- A *list* is a collection of objects each of which can be almost anything

```
> DF <- data.frame(x=1:5, y=letters[1:5])
> DF # data.frame with two columns and 5 rows
  x y
1 1 a
2 2 b
3 3 c
4 4 d
5 5 e
>
> # DF <- data.frame(x=1:10, y=1:7) # illegal because lengths differ
> L <- list(x=1:5, y=1:3, z = DF)
> L # lists are much more flexible!
$x
[1] 1 2 3 4 5

$y
[1] 1 2 3

$z
  x y
1 1 a
```

Data types summary

Key points:

- vector classes include numeric, logical, character, and factors
- vectors can be combined into lists or data.frames
- a data.frame can almost always be thought of as a list of vectors of equal length
- a list is a collection of objects, each of which can be of almost any type

Functions introduced in this section:

`c` combine elements

`as.numeric` convert an object (e.g., a character vector) to numeric

`data.frame` combine object into a data.frame

`ls` list the objects in the workspace

`class` get the class of an object

`str` get the structure of an object

`length` get the number of elements in an object

`mean` calculate the mean of a vector

Exercise 0

- 1 Create a new vector called "test" containing five numbers of your choice [`c()`, `<-`]
- 2 Create a second vector called "students" containing five common names of your choice [`c()`, `<-`]
- 3 Determine the class of "students" and "test" [`class()` or `str()`]
- 4 Create a data frame containing two columns, "students" and "tests" as defined above [`data.frame`]
- 5 Convert "test" to character class, and confirm that you were successful [`as.numeric()`, `<-`, `str()`]

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements**
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Indexing by position or name

Parts of vectors, matrices, data.frames, and lists can be extracted or replaced based on position or name

```
> ## indexing vectors by position
> x <- 101:110 # Create a vector of integers from 101 to 110
> x[c(4, 5)] # extract the fourth and fifth values of x
[1] 104 105
> x[4] <- 1 # change the 4th value to 1
> x # print x
[1] 101 102 103    1 105 106 107 108 109 110
>
> ## indexing vectors by name
> names(x) <- letters[1:10] # give x names
> print(x) #print x
  a    b    c    d    e    f    g    h    i    j
101 102 103    1 105 106 107 108 109 110
> x[c("a", "f")] # extract the values of a and f from x
  a    f
101 106
```

Logical indexing

Elements can also be selected or replaced based on logical (TRUE/FALSE) vectors.

```
> x > 106 # shows which elements of x are > 106
  a      b      c      d      e      f      g      h      i      j
FALSE FALSE FALSE FALSE FALSE FALSE TRUE  TRUE  TRUE  TRUE
> x[x > 106] # selects elements of x where x > 106
  g      h      i      j
107 108 109 110
```

Additional operators useful for logical indexing:

- `==` equal to
- `!=` not equal to
- `>` greater than
- `<` less than
- `>=` greater than or equal to
- `<=` less than or equal to
- `%in%` is included in
- `&` and
- `|` or

```
> x[x > 106 & x <= 108]
```

Indexing matrices

Extraction on matrices operate in two dimensions: first dimension refers to rows, second dimension refers to columns

```
> ## indexing matrices
> # create a matrix
> (M <- cbind(x = 1:5, y = -1:-5, z = c(6, 3, 4, 2, 8)))
      x  y  z
[1,] 1 -1 6
[2,] 2 -2 3
[3,] 3 -3 4
[4,] 4 -4 2
[5,] 5 -5 8
> M[1:3, ] #extract rows 1 through 3, all columns
      x  y  z
[1,] 1 -1 6
[2,] 2 -2 3
[3,] 3 -3 4
> M[c(5, 3, 1), 2:3] # rows 5, 3 and 1, columns 2 and 3
      y  z
[1,] -5 8
[2,] -3 4
[3,] -1 6
> M[M[, 1] %in% 4:2, 2] # second column where first column <=4 & >= 2
[1] -2 -3 -4
```


Indexing lists

Lists can be indexed in the same way as vectors, with the following extension:

```
> # Lists can be indexed with single brackets, similar to vector indexing
> L[c(1, 2)] # the first two elements of L
$x
[1] 1 2 3 4 5

$y
[1] 1 2 3

> L[1] # a list with one element
$x
[1] 1 2 3 4 5

> ## double brackets select the content of a single selected element
> ## effectively taking it out of the list.
> L[[1]] # a vector
[1] 1 2 3 4 5
```

Indexing data.frames

A data.frame can be indexed in the same ways as a matrix, and also the same ways as a list:

```
> DF[c(3, 1, 2), c(1, 2)] # rows 3, 1, and 2, columns 1 and 2
  x y
3 3 c
1 1 a
2 2 b
> DF[[1]] # column 1 as a vector
[1] 1 2 3 4 5
```

There is a subtle but important difference between `[, n]` and `[n]` when indexing data.frames: the first form returns a vector, the second returns a data.frame with one column.

```
> str(DF[1])# a data.frame with one column
'data.frame': 5 obs. of 1 variable:
 $ x: int 1 2 3 4 5
> str(DF[,1])# a vector
int [1:5] 1 2 3 4 5
```

Extraction/replacement summary

Key points:

- elements of objects can be extracted or replaced using the `[]` operator
- objects can be indexed by position, name, or logical (TRUE/FALSE) vectors
- vectors and lists have only one dimension, and hence only one index is used
- matrices and data.frames have two dimensions, and extraction methods for these objects use two indices

Functions introduced in this section:

`[]` extraction operator, used to extract/replace object elements
`names` get the names of an object, usually a vector, list, or data.frame
`print` print an object

Exercise 1

- ❶ Select just the Sepal.Length and Species columns from the *iris* data set (built-in, will be available in your workspace automatically) and save the result to a new data.frame named *iris.2*
- ❷ Calculate the mean of the Sepal.Length column in *iris.2*
- ❸ BONUS (optional): Calculate the mean of Sepal.Length, but only for the setosa species
- ❹ BONUS (optional): Calculate the number of sepal lengths that are more than one standard deviation below the average sepal length

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements**
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

The apply function

The apply function is used to apply a function to the rows or columns of a matrix

```
> M <- matrix(1:20, ncol=4)
> apply(M, 2, mean) ## average across the rows
[1] 3 8 13 18
> apply(M, 2, sum) ## sum the columns
[1] 15 40 65 90
```

The sapply function

It is often useful to apply a function to each element of a vector, list, or data.frame; use the sapply function for this

```
> sapply(DF, class) # get the class of each column in the DF data.frame
      x      y
"integer" "factor"
> sapply(L, length) # get the length of each element in the L list
x y z
5 3 2
> sapply(DF, is.numeric) # check each column of DF to see if it is numeric
      x      y
TRUE FALSE
```

Combining supply and indexing

The `sapply` function can be used in combination with indexing to extract elements that meet certain criteria

- Recall that we can index using logical vectors:

```
> DF[, c(TRUE, FALSE)] # select the first column of DF, but not the second
[1] 1 2 3 4 5
```

```
> ## recall that we can index using logical vectors:
> DF[, c(TRUE, FALSE)] # select the first column of DF, but not the second
[1] 1 2 3 4 5
```

- `sapply()` can be used to generate the logical vector

```
> (DF.which.num <- sapply(DF, is.numeric)) # check which columns of DF are numeric
      x      y
TRUE FALSE
> DF[DF.which.num] # select the numeric columns
  x
1 1
2 2
3 3
4 4
5 5
```


Applying functions summary

Key points:

- R has convenient methods for applying functions to matrices, lists, and data.frames
- other apply-style functions exist, e.g., lapply, tapply, and mapply (see documentation of these functions for details)

Functions introduced in this section:

`matrix` create a matrix (vector with two dimensions)

`apply` apply a function to the rows or columns of a matrix

`sapply` apply a function to the elements of a list

`is.numeric` returns TRUE or FALSE, depending on the type of object

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions**
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Functions

- A function is a collection of commands that takes input(s) and returns output.
 - If you have a specific analysis or transformation you want to do on different data, use a function
 - Functions are defined using the `function()` function
 - Functions can be defined with any number of named arguments
 - Arguments can be of any type (e.g., vectors, data.frames, lists ...)
-

Function return value

- The return value of a function can be:
 - The last object stored in the body of the function
 - Objects explicitly returned with the `return()` function
- Other function output can come from:
 - Calls to `print()`, `message()` or `cat()` in function body
 - Error messages
- Assignment inside the body of a function takes place in a local environment
- Example:

```
> f <- function() { # define function f
+ print("setting x to 1") # print a text string
+ x <- 1} # set x to 1
>
> y <- f() # assign y the value returned by f
[1] "setting x to 1"
>
> y # print y
[1] 1
> x # x in the global is not 1!
  a   b   c   d   e   f   g   h   i   j
101 102 103   1 105 106 107 108 109 110
```

Writing functions example

Goal: write a function that returns the square of it's argument

```
> square <- function (x) { # define function named "square" with argument x
+   return(x*x) # multiple the x argument by itself
+ } # end the function definition
>
> # check to see that the function works
> square(x = 2) # square the value 2
[1] 4
> square(10) # square the value 10
[1] 100
> square(1:5) # square integers 1 through 5
[1] 1 4 9 16 25
```

Debugging basics

Stepping through functions and setting breakpoints

```
> ## Debugging
> # write my.mean function
> my.mean <- function(x, ...) {
+   S <- sum(x, ...)
+   L <- length(na.omit(x))
+   return(S / L)}
> debug(my.mean) # turn debugger on for my.mean function
> # mymean() # step through the function
> undebug(my.mean) # to turn the debugger off
> # insert breakpoints
> my.mean <- function(x, ...) {
+   S <- sum(x, ...)
+   L <- length(na.omit(x))
+   browser() # function will stop here so you can inspect S and L
+   return(S / L)}
>
```

Use `traceback()` to see what went wrong after the fact

```
myModel <- lm(NA~NA)
traceback()
```

Writing functions summary

Key points:

- writing new functions is easy!
- most functions will have a return value, but functions can also print things, write things to file etc.
- functions can be stepped through to facilitate debugging

Functions introduced in this section

`function` defines a new function

`return` used inside a function definition to set the return value

`browser` sets a break point

`debug` turns on the debugging flag of a function so you can step through it

`undebug` turns off the debugging flag

`traceback` shows the error stack (call after an error to see what went wrong)

Exercise 2

- 1 Write a function that takes a `data.frame` as an argument and returns the mean of each numeric column in the data frame. Test your function using the iris data.
- 2 Modify your function so that it returns a list, the first element of which is the means of the numeric variables, the second of which is the counts of the levels of each categorical variable.

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow**
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Control flow

- Basic idea: if some condition is true, do one thing. If false, do something else
 - Carried out in R using `if()` and `else()` statements, which can be nested if necessary
 - Especially useful for checking function arguments and performing different operations depending on function input
-

Control flow examples

Goal: write a function that tells us if a number is positive or negative

```
> ## use branching to return different result depending on the sign of the input
> isPositive <- function(x) { # define function "isPositive"
+   if (x > 0) { # if x is greater than zero, then
+     cat(x, "is positive \n") } # say so!
+   else { # otherwise
+     cat(x, "is negative \n")} # say x is negative
+ } # end function definition
>
> ## test the isPositive() function
> isPositive(10)
10 is positive
> isPositive(-1)
-1 is negative
> isPositive(0)
0 is negative
```

Need to do something different if x equals zero!

Control flow examples

Add a condition to handle $x = 0$

```
> ## add condition to handle the case that x is zero
> isPositive <- function(x) { # define function "isPositive"
+   if (x > 0) { # if x is greater than zero, then
+     cat(x, "is positive \n") } # say so!
+   else if (x == 0) { # otherwise if x is zero
+     cat(x, "is zero \n")} # say so!
+   else { #otherwise
+     cat(x, "is negative \n")} # say x is negative
+ } # end function definition
>
```

Test the new function

```
> isPositive(0) # test the isPositive() function
0 is zero
> isPositive("a") #oops, that will not work!
a is positive
```

We fixed the problem when $x = 0$, but now we need to make sure x is numeric of length one (unless we agree with R that a is positive!)

Control flow examples

Do something reasonable if x is not numeric

```
> ## add condition to handle the case that x is zero
> isPositive <- function(x) { # define function "isPositive"
+   if(!is.numeric(x) | length(x) > 1) {
+     cat("x must be a numeric vector of length one! \n")}
+   else if (x > 0) { # if x is greater than zero, then
+     cat(x, " is positive \n") } # say so!
+   else if (x == 0) { # otherwise if x is zero
+     cat(x, " is zero \n")} # say so!
+   else { #otherwise
+     cat(x, " is negative \n")} # say x is negative
+ } # end function definition
>
> isPositive("a") # test the isPositive() function on character
x must be a numeric vector of length one!
```

Control flow summary

Key points:

- code can be conditionally executed
- conditions can be nested
- conditional execution is often used for argument checking, among other things

Functions introduced in this section

`cat` Concatenates and prints R objects

`if` execute code only if condition is met

`else` used with `if`; code to execute if condition is not met

Exercise 3

- 1 Add argument checking code to return an error if the argument to your function is not a data.frame
- 2 Insert a break point with `browser()` and step through your function

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system**
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

The S3 object class system

R has two major object systems:

- Relatively informal "S3" classes
 - Stricter, more formal "S4" classes
 - We will cover only the S3 system, not the S4 system
 - Basic idea: functions have different methods for different types of objects
-

Object class

The class of an object can be retrieved and modified using the `class()` function:

```
> x <- 1:10
> class(x)
[1] "integer"
> class(x) <- "foo"
> class(x)
[1] "foo"
```

Objects are not limited to a single class, and can have many classes:

```
> class(x) <- c("A", "B")
> class(x)
[1] "A" "B"
```

Function methods

- Functions can have many methods, allowing us to have (e.g.) one `plot()` function that does different things depending on what is being plotted()
- Methods can only be defined for generic functions: `plot`, `print`, `summary`, `mean`, and several others are already generic

```
> # see what methods have been defined for the mean function
> methods(mean)
[1] mean.Date      mean.default    mean.difftime   mean.POSIXct
[5] mean.POSIXlt
> # which functions have methods for data.frames?
> methods(class="data.frame")[1:9]
[1] "[.data.frame"          "[[.data.frame"
[3] "[[<-.data.frame"       "[[<-.data.frame"
[5] "$<-.data.frame"        "aggregate.data.frame"
[7] "anyDuplicated.data.frame" "as.data.frame.data.frame"
[9] "as.list.data.frame"
```

Creating new function methods

To create a new method for a function that is already generic all you have to do is name your function `function.class`

```
> # create a mean() method for objects of class "foo":
> mean.foo <- function(x) { # mean method for "foo" class
+   if(is.numeric(x)) {
+     cat("The average is", mean.default(x))
+     return(invisible(mean.default(x))) #use mean.default for numeric
+   } else
+     cat("x is not numeric \n")} # otherwise say x not numeric
>
> x <- 1:10
> mean(x)
[1] 5.5
> class(x) <- "foo"
> mean(x)
The average is 5.5>
> x <- as.character(x)
> class(x) <- "foo"
> mean(x)
x is not numeric
```

Creating generic functions

S3 generics are most often used for print, summary, and plot methods, but sometimes you may want to create a new generic function

```
> # create a generic disp() function
> disp <- function(x, ...) {
+   UseMethod("disp")
+ }
>
> # create a disp method for class "matrix"
> disp.matrix <- function(x) {
+   print(round(x, digits=2))
+ }
>
> # test it out
> disp(matrix(runif(10), ncol=2))
      [,1] [,2]
[1,] 0.45 0.79
[2,] 0.96 0.36
[3,] 0.14 0.06
[4,] 0.78 0.57
[5,] 0.80 0.66
```

S3 classes summary

Key points:

- there are several class systems in R, of which S3 is the oldest and simplest
- objects have *class* and functions have corresponding *methods*
- the class of an object can be set by simple assignment
- S3 generic functions all contain `UseMethod("x")` in the body, where `x` is the name of the function
- new methods for existing generic functions can be written simply by defining a new function with a special naming scheme: the name of the function followed by dot followed by the name of the class

Functions introduced in this section

`plot` creates a graphical display, the type of which depends on the class of the object being plotted

`methods` lists the methods defined for a function or class

`UseMethod` the body of a generic function

`invisible` returns an object but does not print it

Exercise 4

- 1 Modify your function so that it also returns the standard deviations of the numeric variables
- 2 Modify your function so that it returns a list of class "statsum"
- 3 Write a print method for the statsum class

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you**
- 9 Additional resources
- 10 Loops (supplimental)

Gotcha's

- There are an unfortunately large number of surprises in R programming
 - Some of these "gotcha's" are common problems in other languages, many are unique to R
 - We will only cover a few – for a more comprehensive discussion please see http://www.burns-stat.com/pages/Tutor/R_inferno.pdf
-

Floating point comparison

Floating point arithmetic is not exact:

```
> .1 == .3/3  
[1] FALSE
```

Solution: use `all.equal()`:

```
> all.equal(.1, .3/3)  
[1] TRUE
```

Missing values

R does not exclude missing values by default – a single missing value in a vector means that many things are unknown:

```
> x <- c(1:10, NA, 12:20)
> c(mean(x), sd(x), median(x), min(x), sd(x))
[1] NA NA NA NA NA
```

NA is not equal to anything, not even NA

```
> NA == NA
[1] NA
```

Solutions: use `na.rm = TRUE` option when calculating, and `is.na` to test for missing

Automatic type conversion

Automatic type conversion happens a lot which is often useful, but makes it easy to miss mistakes

```
> # combining values coerces them to the most general type
> (x <- c(TRUE, FALSE, 1, 2, "a", "b"))
[1] "TRUE" "FALSE" "1" "2" "a" "b"
> str(x)
chr [1:6] "TRUE" "FALSE" "1" "2" "a" "b"
>
> # comparisons convert arguments to most general type
> 1 > "a"
[1] FALSE
```

Maybe this is what you expect... I would like to at least get a warning!

Optional argument inconsistencies

Functions you might expect to work similarly don't always:

```
> mean(1, 2, 3, 4, 5)*5  
[1] 5  
> sum(1, 2, 3, 4, 5)  
[1] 15
```

Why are these different?!?

```
> args(mean)  
function (x, ...)  
NULL  
> args(sum)  
function (... , na.rm = FALSE)  
NULL
```

Ouch. That is not nice at all!

Trouble with Factors

Factors sometimes behave as numbers, and sometimes as characters, which can be confusing!

```
> (x <- factor(c(5, 5, 6, 6), levels = c(6, 5)))
[1] 5 5 6 6
Levels: 6 5
>
> str(x)
Factor w/ 2 levels "6","5": 2 2 1 1
>
> as.character(x)
[1] "5" "5" "6" "6"
> # here is where people sometimes get lost...
> as.numeric(x)
[1] 2 2 1 1
> # you probably want
> as.numeric(as.character(x))
[1] 5 5 6 6
```

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources**
- 10 Loops (supplimental)

Additional reading and resources

- S3 system overview:
<https://github.com/hadley/devtools/wiki/S3>
- S4 system overview:
<https://github.com/hadley/devtools/wiki/S4>
- R documentation: <http://cran.r-project.org/manuals.html>
- Collection of R tutorials:
<http://cran.r-project.org/other-docs.html>
- R for Programmers (by Norman Matloff, UC-Davis)

<http://heather.cs.ucdavis.edu/~matloff/R/RProg.pdf>

- Calling C and Fortran from R (by Charles Geyer, UMinn)

<http://www.stat.umn.edu/~charlie/rc/>

- State of the Art in Parallel Computing with R (Schmidberger et al.)

<http://www.jstatsoft.org/v31/i01/paper>

- Institute for Quantitative Social Science: <http://iq.harvard.edu>
- Research technology consulting:
<http://projects.iq.harvard.edu/rtc>

Feedback

- Help Us Make This Workshop Better!
- Please take a moment to fill out a very short feedback form
- These workshops exist for you – tell us what you need!
- <http://tinyurl.com/RprogrammingFeedback>

Topic

- 1 Workshop overview and materials
- 2 Data types
- 3 Extracting and replacing object elements
- 4 Applying functions to list elements
- 5 Writing functions
- 6 Control flow
- 7 The S3 object class system
- 8 Things that may surprise you
- 9 Additional resources
- 10 Loops (supplimental)

Looping

- A loop is a collection of commands that are run over and over again.
 - A for loop runs the code a fixed number of times, or on a fixed set of objects.
 - A while loop runs the code until a condition is met.
 - If you're typing the same commands over and over again, you might want to use a loop!
-

Looping: for-loop examples

For each value in a vector, print the number and its square

```
> # For-loop example
> for (num in seq(-5,5)) {# for each number in [-5, 5]
+   cat(num, "squared is", num^2, "\n") # print the number
+ }
```

-5 squared is 25
-4 squared is 16
-3 squared is 9
-2 squared is 4
-1 squared is 1
0 squared is 0
1 squared is 1
2 squared is 4
3 squared is 9
4 squared is 16
5 squared is 25

Looping: while-loop example

Goal: simulate rolling two dice until we roll two sixes

```
> ## While-loop example: rolling dice
> set.seed(15) # allows reproducible sample() results
> dice <- seq(1,6) # set dice = [1 2 3 4 5 6]
> roll <- 0 # set roll = 0
> while (roll < 12) {
+   roll <- sample(dice,1) + sample(dice,1) # calculate sum of two rolls
+   cat("We rolled a ", roll, "\n") # print the result
+ } # end the loop
We rolled a 6
We rolled a 10
We rolled a 9
We rolled a 7
We rolled a 10
We rolled a 5
We rolled a 9
We rolled a 12
```

Using loops to fill in lists

Often you will want to store the results from a loop. You can create an object to hold the results generated in the loop and fill in the values using indexing

```
> ## save calculations done in a loop
> Result <- list() # create an object to store the results
> for (i in 1:5) {# for each i in [1, 5]
+   Result[[i]] <- 1:i ## assign the sequence 1 to i to Result
+ }
> Result # print Result
[[1]]
[1] 1

[[2]]
[1] 1 2

[[3]]
[1] 1 2 3

[[4]]
[1] 1 2 3 4

[[5]]
[1] 1 2 3 4 5
```

Word of caution: don't overuse loops!

Most operations in R are vectorized – This makes loops unnecessary in many cases

- Use vector arithmetic instead of loops:

```
> x <- c() # create vector x
> for(i in 1:5) x[i] <- i+i # double a vector using a loop
> print(x) # print the result
[1] 2 4 6 8 10
>
> 1:5 + 1:5 #double a vector without a loop
[1] 2 4 6 8 10
> 1:5 + 5 # shorter vectors are recycled
[1] 6 7 8 9 10
```

- Use paste instead of loops:

```
> ## Earlier we said
> ## for (num in seq(-5,5)) {# for each number in [-5, 5]
> ##   cat(num, "squared is", num^2, "\n") # print the number
> ## }
> ## a better way:
> paste(1:5, "squared = ", (1:5)^2)
[1] "1 squared = 1" "2 squared = 4" "3 squared = 9"
[4] "4 squared = 16" "5 squared = 25"
```

Exercise 5

- 1 use a loop to get the `class()` of each column in the iris data set
- 2 use the results from step 1 to select the numeric columns
- 3 use a loop to calculate the mean of each numeric column in the iris data