

Regression Models in R



The Institute
for Quantitative Social Science
at Harvard University

Outline

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors
- 4 Regression with binary outcomes
- 5 Multilevel Modeling
- 6 Multiple imputation
- 7 Wrap-up

Topic

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors
- 4 Regression with binary outcomes
- 5 Multilevel Modeling
- 6 Multiple imputation
- 7 Wrap-up

Workshop description

- This is an intermediate/advanced R course
- Appropriate for those with basic knowledge of R
- This is not a statistics course!
- Learning objectives:
 - Learn the R formula interface
 - Specify factor contrasts to test specific hypotheses
 - Perform model comparisons
 - Run and interpret variety of regression models in R
 - Create and use imputed data sets in regression models

Materials and Setup

- Lab computer users:
 USERNAME dataclass
 PASSWORD dataclass
- Download materials from `http:j.mp/r-stats`
- Extract materials from `RStatistics.zip` (on lab machines *right-click -> WinZip -> Extract to here*) and move to your desktop.

Launch RStudio

- Open the RStudio program from the Windows start menu
- Open up today's R script
 - In RStudio, Go to **File => Open Script**
 - Locate and open the Rstatistics.R script in the Rstatistics folder on your desktop
- Go to **Tools => Set working directory => To source file location** (more on the working directory later)
- I encourage you to add your own notes to this file!

Set working directory

It is often helpful to start your R session by setting your working directory so you don't have to type the full path names to your data and other files

```
> # set the working directory
> # setwd("~/Desktop/Rstatistics")
> # setwd("C:/Users/dataclass/Desktop/Rstatistics")
>
```

You might also start by listing the files in your working directory

```
> getwd() # where am I?
[1] "/home/izahn/Documents/Work/IQSS/Classes/IQSS_Stats_Workshops/R/Rstatistics"
> list.files("dataSets") # files in the dataSets folder
[1] "Exam.rds"          "NatHealth2008MI"    "NatHealth2011.rds"
[4] "states.dta"         "states.rds"
```

Load the states data

The *states.dta* data comes from <http://anawida.de/teach/SS14/anawida/4.linReg/data/states.dta.txt> and appears to have originally appeared in *Statistics with Stata* by Lawrence C. Hamilton.

```
> # read the states data
> states.data <- readRDS("dataSets/states.rds")
> #get labels
> states.info <- data.frame(attributes(states.data)[c("names", "var.labels")])
> #look at last few labels
> tail(states.info, 8)
```

	names	var.labels
14	csat	Mean composite SAT score
15	vsat	Mean verbal SAT score
16	msat	Mean math SAT score
17	percent	% HS graduates taking SAT
18	expense	Per pupil expenditures prim&sec
19	income	Median household income, \$1,000
20	high	% adults HS diploma
21	college	% adults college degree

Topic

- 1 Introduction
- 2 Linear regression**
- 3 Interactions and factors
- 4 Regression with binary outcomes
- 5 Multilevel Modeling
- 6 Multiple imputation
- 7 Wrap-up

Examine the data before fitting models


Start by examining the data to check for problems.

```
> # summary of expense and csat columns, all rows
> summary(states.data[, c("expense", "csat")])
      expense      csat
Min.   :2960   Min.   : 832
1st Qu.:4352   1st Qu.: 888
Median :5000   Median : 926
Mean   :5236   Mean    : 944
3rd Qu.:5794   3rd Qu.: 997
Max.   :9259   Max.    :1093
> # correlation between expense and csat
> cor(states.data[, c("expense", "csat")])[1, 2]
[1] -0.466
```

Plot the data before fitting models

Plot the data to look for multivariate outliers, non-linear relationships etc.

```
# scatter plot of expense vs csat  
plot(states.data[ , c("expense", "csat")])
```



Linear regression example

- Linear regression models can be fit with the `lm()` function
- For example, we can use `lm` to predict SAT scores based on per-pupil expenditures:

```
> # Fit our regression model
> sat.mod <- lm(csat ~ expense, # regression formula
+              data=states.data) # data set
> # Summarize and print the results
> summary(sat.mod) # show regression coefficients table
```

Call:

```
lm(formula = csat ~ expense, data = states.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-131.81	-38.08	5.61	37.85	136.50

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1060.73244	32.70090	32.44	< 2e-16
expense	-0.02228	0.00604	-3.69	0.00056

Residual standard error: 59.8 on 49 degrees of freedom

Multiple R-squared: 0.217 Adjusted R-squared: 0.201

Why is the association between expense and SAT scores *negative*?

Many people find it surprising that the per-capita expenditure on students is negatively related to SAT scores. The beauty of multiple regression is that we can try to pull these apart. What would the association between expense and SAT scores be if there were no difference among the states in the percentage of students taking the SAT?

```
> summary(lm(csats ~ expense + percent, data = states.data))
```

Call:

```
lm(formula = csats ~ expense + percent, data = states.data)
```

Residuals:

Min	1Q	Median	3Q	Max
-62.92	-24.32	1.74	15.50	75.62

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	989.8074	18.3958	53.81	< 2e-16
expense	0.0086	0.0042	2.05	0.046
percent	-2.5377	0.2249	-11.28	4.2e-15

The lm class and methods

OK, we fit our model. Now what?

- Examine the model object:

```
> class(sat.mod)
[1] "lm"
> names(sat.mod)
[1] "coefficients" "residuals"      "effects"        "rank"
[5] "fitted.values" "assign"         "qr"            "df.residual"
[9] "xlevels"      "call"          "terms"         "model"
> methods(class = class(sat.mod))[1:9]
[1] "add1.lm"          "alias.lm"        "anova.lm"
[4] "case.names.lm"    "confint.lm"      "cooks.distance.lm"
[7] "deviance.lm"      "dfbeta.lm"       "dfbetas.lm"
```


- Use function methods to get more information about the fit

```
> confint(sat.mod)
                2.5 %      97.5 %
(Intercept) 995.0175 1126.4474
expense      -0.0344  -0.0101
> # hist(residuals(sat.mod))
>
```

Linear Regression Assumptions

- Ordinary least squares regression relies on several assumptions, including that the residuals are normally distributed and homoscedastic, the errors are independent and the relationships are linear.
- Investigate these assumptions visually by plotting your model:

```
par(mar = c(4, 4, 2, 2), mfrow = c(1, 2)) #optional
plot(sat.mod, which = c(1, 2)) # "which" argument optional
```



Comparing models

Do congressional voting patterns predict SAT scores over and above expense?
Fit two models and compare them:

```
> # fit another model, adding house and senate as predictors
> sat.voting.mod <- lm(csat ~ expense + house + senate,
+                      data = na.omit(states.data))
> sat.mod <- update(sat.mod, data=na.omit(states.data))
> # compare using the anova() function
> anova(sat.mod, sat.voting.mod)
```

Analysis of Variance Table

Model 1: csat ~ expense

Model 2: csat ~ expense + house + senate

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	46	169050				
2	44	149284	2	19766	2.91	0.065

```
> coef(summary(sat.voting.mod))
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1082.9344	38.63381	28.03	1.07e-29
expense	-0.0187	0.00969	-1.93	6.00e-02
house	-1.4424	0.60048	-2.40	2.06e-02
senate	0.4982	0.51356	0.97	3.37e-01

Exercise 0: least squares regression

Use the *states.rds* data set. Fit a model predicting energy consumed per capita (energy) from the percentage of residents living in metropolitan areas (metro). Be sure to

- 1 Examine/plot the data before fitting the model
- 2 Print and interpret the model `summary`
- 3 plot the model to look for deviations from modeling assumptions

Select one or more additional predictors to add to your model and repeat steps 1-3. Is this model significantly better than the model with *metro* as the only predictor?

Topic

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors**
- 4 Regression with binary outcomes
- 5 Multilevel Modeling
- 6 Multiple imputation
- 7 Wrap-up

Modeling interactions

Interactions allow us assess the extent to which the association between one predictor and the outcome depends on a second predictor. For example: Does the association between expense and SAT scores depend on the median income in the state?

```
> #Add the interaction to the model
> sat.expense.by.percent <- lm(csat ~ expense*income,
+                               data=states.data)
> #Show the results
> coef(summary(sat.expense.by.percent)) # show regression coefficients table
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1380.36423	172.086252	8.02	0.000000000237
expense	-0.06384	0.032701	-1.95	0.056878369245
income	-10.49785	4.991463	-2.10	0.040832525071
expense:income	0.00138	0.000864	1.60	0.115539488253

Regression with categorical predictors

Let's try to predict SAT scores from region, a categorical variable. Note that you must make sure R does not think your categorical variable is numeric.

```
> # make sure R knows region is categorical
> str(states.data$region)
Factor w/ 4 levels "West","N. East",...: 3 1 1 3 1 1 2 3 NA 3 ...
> states.data$region <- factor(states.data$region)
> #Add region to the model
> sat.region <- lm(csat ~ region,
+                 data=states.data)
> #Show the results
> coef(summary(sat.region)) # show regression coefficients table
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	946.3	14.8	63.958	1.35e-46
regionN. East	-56.8	23.1	-2.453	1.80e-02
regionSouth	-16.3	19.9	-0.819	4.17e-01
regionMidwest	63.8	21.4	2.986	4.51e-03

```
> anova(sat.region) # show ANOVA table
Analysis of Variance Table

Response: csat
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
region	3	82049	27350	9.61	0.000049
Residuals	46	130912	2846		

Setting factor reference groups and contrasts

In the previous example we use the default contrasts for region. The default in R is treatment contrasts, with the first level as the reference. We can change the reference group or use another coding scheme using the `C` function.

```
> # print default contrasts
> contrasts(states.data$region)
```

	N. East	South	Midwest
West	0	0	0
N. East	1	0	0
South	0	1	0
Midwest	0	0	1

```
> # change the reference group
> coef(summary(lm(csat ~ C(region, base=4),
+               data=states.data)))
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1010.1	15.4	65.59	4.30e-47
C(region, base = 4)1	-63.8	21.4	-2.99	4.51e-03
C(region, base = 4)2	-120.5	23.5	-5.12	5.80e-06
C(region, base = 4)3	-80.1	20.4	-3.93	2.83e-04

```
> # change the coding scheme
> coef(summary(lm(csat ~ C(region, contr.helmert),
+               data=states.data)))
```

	Estimate	Std. Error	t value	Pr(> t)
--	----------	------------	---------	----------

Exercise 1: interactions and factors

Use the states data set.

- 1 Add on to the regression equation that you created in exercise 1 by generating an interaction term and testing the interaction.
- 2 Try adding a categorical variable to your regression (remember, it will need to be dummy coded). You could use region or, or generate a new categorical variable from one of the continuous variables in the dataset.
- 3 Are there significant differences across the four regions?

Topic

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors
- 4 Regression with binary outcomes**
- 5 Multilevel Modeling
- 6 Multiple imputation
- 7 Wrap-up

Logistic regression

This far we have used the `lm` function to fit our regression models. `lm` is great, but limited—in particular it only fits models for continuous dependent variables. For categorical dependent variables we can use the `glm()` function. For these models we will use a different dataset, drawn from the National Health Interview Survey. From the CDC website:

The National Health Interview Survey (NHIS) has monitored the health of the nation since 1957. NHIS data on a broad range of health topics are collected through personal household interviews. For over 50 years, the U.S. Census Bureau has been the data collection agent for the National Health Interview Survey. Survey results have been instrumental in providing data to track health status, health care access, and progress toward achieving national health objectives.

Load the National Health Interview Survey data:

```
> NH11 <- readRDS("dataSets/NatHealth2011.rds")  
> labs <- attributes(NH11)$labels  
>
```


Logistic regression example

Let's predict the probability of being diagnosed with hypertension based on age, sex, sleep, and bmi

```
> str(NH11$hypev) # check stucture of hypev
Factor w/ 5 levels "1 Yes","2 No",...: 2 2 1 2 2 1 2 2 1 2 ...
> levels(NH11$hypev) # check levels of hypev
[1] "1 Yes"          "2 No"           "7 Refused"
[4] "8 Not ascertained" "9 Don't know"
> # collapse all missing values to NA
> NH11$hypev <- factor(NH11$hypev, levels=c("2 No", "1 Yes"))
> # run our regression model
> hyp.out <- glm(hypev~age_p+sex+sleep+bmi,
+               data=NH11, family="binomial")
> coef(summary(hyp.out))
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-4.26947	0.056495	-75.57	0.00e+00
age_p	0.06070	0.000823	73.78	0.00e+00
sex2 Female	-0.14403	0.026798	-5.37	7.68e-08
sleep	-0.00704	0.001640	-4.29	1.78e-05
bmi	0.01857	0.000951	19.53	6.49e-85

Logistic regression coefficients

Generalized linear models use link functions, so raw coefficients are difficult to interpret. For example, the age coefficient of .06 in the previous model tells us that for every one unit increase in age, the log odds of hypertension diagnosis increases by 0.06. Since most of us are not used to thinking in log odds this is not too helpful!

One solution is to transform the coefficients to make them easier to interpret

```
> hyp.out.tab <- coef(summary(hyp.out))
> hyp.out.tab[, "Estimate"] <- exp(coef(hyp.out))
> hyp.out.tab
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.014	0.056495	-75.57	0.00e+00
age_p	1.063	0.000823	73.78	0.00e+00
sex2 Female	0.866	0.026798	-5.37	7.68e-08
sleep	0.993	0.001640	-4.29	1.78e-05
bmi	1.019	0.000951	19.53	6.49e-85

Generating predicted values

In addition to transforming the log-odds produced by `glm` to odds, we can use the `predict()` function to make direct statements about the predictors in our model. For example, we can ask "How much more likely is a 63 year old female to have hypertension compared to a 33 year old female?".

```
> # Create a dataset with predictors set at desired levels
> predDat <- with(NH11,
+               expand.grid(age_p = c(33, 63),
+                           sex = "2 Female",
+                           bmi = mean(bmi, na.rm = TRUE),
+                           sleep = mean(sleep, na.rm = TRUE)))
> # predict hypertension at those levels
> cbind(predDat, predict(hyp.out, type = "response",
+                       se.fit = TRUE, interval="confidence",
+                       newdata = predDat))
```

	age_p	sex	bmi	sleep	fit	se.fit	residual.scale
1	33	2 Female	29.9	7.86	0.129	0.00285	1
2	63	2 Female	29.9	7.86	0.478	0.00482	1


This tells us that a 33 year old female has a 13% probability of having been diagnosed with hypertension, while a 63 year old female has a 48% probability of having been diagnosed.

Packages for computing and graphing predicted values


Instead of doing all this ourselves, we can use the effects package to compute quantities of interest for us (cf. the Zelig package).

```
library(effects)
plot(allEffects(hyp.out))
```


age_p effect plot




sex effect plot



sleep effect plot



bmi effect plot



Exercise 2: logistic regression

Use the NH11 data set.

- 1 Use glm to conduct a logistic regression to predict ever worked (everwrk) using age (age_p) and marital status (r_{maritl}).
- 2 Predict the probability of working for each level of marital status.

Note that the data is not perfectly clean and ready to be modeled. You will need to clean up at least some of the variables before fitting the model.

Topic

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors
- 4 Regression with binary outcomes
- 5 Multilevel Modeling**
- 6 Multiple imputation
- 7 Wrap-up

Multilevel modeling overview

- Multi-level (AKA hierarchical) models are a type of mixed-effects models
- Used to model variation due to group membership where the goal is to generalize to a population of groups
- Can model different intercepts and/or slopes for each group
- Mixed-effects models include two types of predictors: fixed-effects and random effects

Fixed-effects observed levels are of direct interest (.e.g, sex, political party. . .)

Random-effects observed levels not of direct interest: goal is to make inferences to a population represented by observed levels

- In R the lme4 package is the most popular for mixed effects models
 - Use the `lmer` function for liner mixed models, `glmer` for generalized mixed models

```
> library(lme4)
>
```

The Exam data

The Exam data set contains exam scores of 4,059 students from 65 schools in Inner London. The variable names are as follows:

- school** School ID - a factor.
- normexam** Normalized exam score.
- schgend** School gender - a factor. Levels are 'mixed', 'boys', and 'girls'.
- schavg** School average of intake score.
- vr** Student level Verbal Reasoning (VR) score band at intake - a factor. Levels are 'bottom 25%', 'mid 50%', and 'top 25%'.
- intake** Band of student's intake score - a factor. Levels are 'bottom 25%', 'mid 50%' and 'top 25%'./
- standLRT** Standardised LR test score.
- sex** Sex of the student - levels are 'F' and 'M'.
- type** School type - levels are 'Mxd' and 'Sngl'.
- student** Student id (within school) - a factor

```
> Exam <- readRDS("dataSets/Exam.rds")  
>
```


The null model and ICC

As a preliminary step it is often useful to partition the variance in the dependent variable into the various levels. This can be accomplished by running a null model (i.e., a model with a random effects grouping structure, but no fixed-effects predictors).

```
> # null model, grouping by school but not fixed effects.
> Norm1 <- lmer(normexam ~ 1 + (1|school),
+               data=Exam, REML = FALSE)
> summary(Norm1)
```

```
Linear mixed model fit by maximum likelihood ['lmerMod']
Formula: normexam ~ 1 + (1 | school)
Data: Exam
```

AIC	BIC	logLik	deviance	df.resid
10826	10844	-5410	10820	3984

Scaled residuals:

Min	1Q	Median	3Q	Max
-3.902	-0.646	0.003	0.698	3.636

Random effects:

Groups	Name	Variance	Std.Dev.
school	(Intercept)	0.169	0.412
Residual		0.848	0.921

Adding fixed-effects predictors

Predict exam scores from student's standardized tests scores

```
> Norm2 <- lmer(normexam ~ standLRT + (1|school),
+               data=Exam,
+               REML = FALSE)
> summary(Norm2)
```

Linear mixed model fit by maximum likelihood ['lmerMod']
 Formula: normexam ~ standLRT + (1 | school)
 Data: Exam

	AIC	BIC	logLik	deviance	df.resid
	9143	9168	-4568	9135	3954

Scaled residuals:

	Min	1Q	Median	3Q	Max
	-3.700	-0.625	0.024	0.678	3.262

Random effects:

Groups	Name	Variance	Std.Dev.
school	(Intercept)	0.0919	0.303
Residual		0.5670	0.753

Number of obs: 3958, groups: school, 65

Fixed effects:

Multiple degree of freedom comparisons

As with `lm` and `glm` models, you can compare the two `lmer` models using the `anova` function.

```
> anova(Norm1, Norm2)
Data: Exam
Models:
Norm1: normexam ~ 1 + (1 | school)
Norm2: normexam ~ standLRT + (1 | school)
      Df    AIC    BIC logLik deviance Chisq Chi Df Pr(>Chisq)
Norm1  3 10826 10844  -5410    10820
Norm2  4  9143  9169  -4568     9135  1684    1    <2e-16
```

Random slopes

Add a random effect of students' standardized test scores as well. Now in addition to estimating the distribution of intercepts across schools, we also estimate the distribution of the slope of exam on standardized test.

```
> Norm3 <- lmer(normexam~standLRT + (standLRT|school), data=Exam,
+               REML = FALSE)
```

```
> summary(Norm3)
```

Linear mixed model fit by maximum likelihood ['lmerMod']

Formula: normexam ~ standLRT + (standLRT | school)

Data: Exam

AIC	BIC	logLik	deviance	df.resid
9108	9146	-4548	9096	3952

Scaled residuals:

Min	1Q	Median	3Q	Max
-3.813	-0.634	0.033	0.673	3.452

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
school	(Intercept)	0.0899	0.300	
	standLRT	0.0141	0.119	0.51
Residual		0.5552	0.745	

Number of obs: 3958, groups: school, 65

Test the significance of the random slope

To test the significance of a random slope just compare models with and without the random slope term

```
> anova(Norm2, Norm3)
Data: Exam
Models:
Norm2: normexam ~ standLRT + (1 | school)
Norm3: normexam ~ standLRT + (standLRT | school)
      Df  AIC  BIC logLik deviance Chisq Chi Df    Pr(>Chisq)
Norm2  4 9143 9169  -4568     9135
Norm3  6 9108 9146  -4548     9096    39      2 0.0000000035
```

Exercise 3: multilevel modeling

Use the dataset, bh1996:

```
data(bh1996, package="multilevel")
```

From the data documentation:

Variables are Cohesion (COHES), Leadership Climate (LEAD), Well-Being (WBEING) and Work Hours (HRS). Each of these variables has two variants - a group mean version that replicates each group mean for every individual, and a within-group version where the group mean is subtracted from each individual response. The group mean version is designated with a G. (e.g., G.HRS), and the within-group version is designated with a W. (e.g., W.HRS).

- 1 Create a null model predicting wellbeing ("WBEING")
- 2 Calculate the ICC for your null model
- 3 Run a second multi-level model that adds two individual-level predictors, average number of hours worked ("HRS") and leadership skills ("LEAD") to the model and interpret your output.
- 4 Now, add a random effect of average number of hours worked ("HRS") to the model and interpret your output. Test the significance of this random term.
- 5 Finally, add a group-level term, workplace cohesion ("G_COHES") to the

Topic

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors
- 4 Regression with binary outcomes
- 5 Multilevel Modeling
- 6 Multiple imputation**
- 7 Wrap-up

Multiple imputation

- Majority of datasets contain missing data
- Produces a variety of problems and limitations to data analysis
- Multiple imputation (MI) generates multiple, complete datasets that contain estimations of missing data points

Multiple imputation

Earlier we wanted to compare a model predicting bmi from demographic variables to a model including demographics and substantive predictors. We omitted missing data so that we could fit both models to the same data. That is a common practice, but it has many problems (which we unfortunately don't have time to discuss in detail). A popular solution is to use multiple imputation to fill in the missing values with reasonable placeholders.

MI is typically thought of as involving three steps:

- Selection of imputation model
- Generation of imputed datasets
- Combining results across imputed datasets

There are a number of packages for doing this in R: we will use the Amelia package because it is powerful, fast, and easy to use. You can refer to the Amelia documentation for more information about its imputation procedures: <http://r.iq.harvard.edu/docs/amelia/amelia.pdf>

Creating imputed data sets







We're going to create several datasets to look at a model predicting the number of days of work missed/year (wkdayr)

```
> # load the Amelia package
> library(Amelia)
> # help(package="Amelia")
> # load a smaller version of NH
> NH08.mi <- readRDS("dataSets/NatHealth2008MI")
> # generate five imputed data sets
> amelia.log <- capture.output( # suppress amelia's chattiness
+   NatHealth.MI <- amelia(NH08.mi,
+                           m=5,
+                           idvars=c("id")))
>
```

Checking imputed values

Compare imputed values to observed values

```
plot(NatHealth.MI, which.vars=9:12)
```




Checking imputed values: overimputation

Overimputation strategy:

- Treat every observed value as if it was missing
- Impute many values for that observed value
- Examine the correspondence between imputed and observed values

```
overimpute(NatHealth.MI, var="sleep")
```



Using imputed data sets in regression models

Zelig makes it very easy to use imputed data sets – just point to the list of imputed data sets in the data argument

```
> library(Zelig)
> nhImp.out <- zelig(wkdayr ~ cigsdays + modmin + sleep, model = "ls",
+                   data = NatHealth.MI$imputations, cite = FALSE)
>
> coef(summary(nhImp.out))
```

	Value	Std. Error	t-stat	p-value
(Intercept)	-6.0745	5.6803	-1.07	0.2963
cigsdays	0.0252	0.1326	0.19	0.8496
modmin	-0.0318	0.0227	-1.40	0.1610
sleep	1.7233	0.8624	2.00	0.0669

For separate results, use `print(summary(x), subset = i:j)`.

Exercise 2: multiple imputation

- 1 Using Amelia, generate 5 imputed versions of the Exam dataset. Make sure you tell Amelia which variables are nominal, and that school is the id variable.
- 2 Create plots that compare imputed values to observed values
- 3 Overimpute the variable "schavg"

Topic

- 1 Introduction
- 2 Linear regression
- 3 Interactions and factors
- 4 Regression with binary outcomes
- 5 Multilevel Modeling
- 6 Multiple imputation
- 7 Wrap-up**

Help us make this workshop better!

- Please take a moment to fill out a very short feedback form
- These workshops exist for you – tell us what you need!
- <http://tinyurl.com/RstatisticsFeedback>

Additional resources

- IQSS workshops:
http://projects.iq.harvard.edu/rtc/filter_by/workshops
- IQSS statistical consulting: <http://rtc.iq.harvard.edu>
- Zelig
 - Website: <http://gking.harvard.edu/zelig>
 - Documentation: <http://r.iq.harvard.edu/docs/zelig.pdf>
- Ameila
 - Website: <http://gking.harvard.edu/Amelia/>
 - Documetation: <http://r.iq.harvard.edu/docs/amelia/amelia.pdf>